

Geometric Deep Learning for Structured and Relational AI in Gemantria

Introduction to Geometric Deep Learning (Geometric AI)

Geometric Deep Learning (GDL) is an emerging paradigm that extends deep neural networks to data with **non-Euclidean structure** – such as graphs, manifolds, and symmetric spaces – by leveraging the underlying geometry and symmetry in the data[1][2]. In contrast to traditional deep learning on grids (images, sequences), geometric approaches incorporate **inductive biases** like graph connectivity, spatial transformations, or group symmetries directly into neural network architectures. According to Michael Bronstein (who coined the term "geometric AI"), many high-dimensional tasks are tractable only because their data has *low-dimensional structure* and symmetry; by "exposing these regularities through unified geometric principles," we can build more powerful and **explainable** models[3]. GDL provides a common mathematical blueprint to derive architectures such as CNNs, graph neural networks (GNNs), and even Transformers from first principles of symmetry and invariance[2]. In practical terms, this means designing neural layers that respect the *structure* of the data domain (e.g. the relationships in a graph or the rotations of a molecule) so that the model's computations **commute with** (or are invariant to) transformations that we know should not change the data's meaning.

Core idea: By building the geometry of the domain into our models, we obtain **structured and relational reasoning systems** that generalize better and are easier to interpret. For example, a convolutional network *hard-codes* translation symmetry on image grids, greatly reducing parameters while preserving important features regardless of position[4]. Likewise, a graph neural network respects the arbitrary ordering of nodes by using permutation-invariant message-passing, so that relabeling nodes doesn't change the output[5]. These geometric priors act as *relational inductive biases*, guiding the model to learn in ways that align with known structure in the data[6]. The Gemantria project, with its emphasis on **concept graphs**, **temporal patterns**, and **semantic relationships** in biblical texts, stands to benefit greatly from such techniques. By integrating geometric deep learning into Gemantria's stack, we can enable more **structured reasoning** (through graph-based neural networks), improve **forecasting** by accounting for relational patterns over time, and enhance **semantic enrichment** by embedding and propagating knowledge in concept networks. In the following sections, we break down key concepts of geometric deep learning – focusing on graphs, manifolds, group equivariance, and relational biases – and provide concrete guidance on applying them to Gemantria's use cases.

Geometric Deep Learning on Graphs: Graph Neural Networks and Message Passing

Graphs are a natural way to represent relational knowledge: nodes denote entities or concepts and edges denote relationships. **Graph Neural Networks (GNNs)** are the workhorse of geometric deep learning on graphs, generalizing neural networks to graph-structured data[1]. The fundamental principle is **message-passing** (also called neighborhood aggregation): each node in the graph aggregates information from its neighbors to update its own representation. By stacking such layers, a GNN allows information to propagate along edges, effectively learning *functions on the graph* that

depend on the structure. Crucially, this is done in a way that is **permutation-invariant** to node ordering – any relabeling of nodes yields the same aggregated results[5]. This ensures the network respects the graph's symmetry (the fact that nodes have no implicit sequence) and focuses only on the relational pattern.

How GNNs work: In a typical message-passing layer, each node i receives “messages” from its adjacent nodes $j \in \mathcal{N}(i)$, transforms or weights those messages, and combines them (often by summation or averaging) to update node i 's feature vector. For instance, a simple Graph Convolutional Network (GCN) update rule is:

$$\text{h}_i^{\text{(new)}} = \sigma(W \sum_{j \in \mathcal{N}(i)} \cup \{j\} \frac{1}{\sqrt{d_i d_j}} \text{h}_j^{\text{(old)}}),$$

where h_j are neighbor features, W is a trainable weight matrix, d_i, d_j are degree-based normalizers, and σ is a nonlinearity. This **localized convolution** on the graph mimics a traditional CNN filter but on an irregular neighborhood[7]. More advanced GNN variants use different aggregation functions or attention mechanisms to weigh neighbors:

- **Graph Attention Networks (GAT):** assign an attention weight to each edge (i,j) based on the interacting features of node i and j , so that important neighbors influence i 's representation more. This can improve performance and offers interpretability by indicating which relations were most influential.
- **GraphSAGE:** learns an aggregation function (mean, LSTM-based, etc.) that can generalize to unseen nodes. It enables inductive reasoning on graphs by sampling neighbors and aggregating, rather than requiring a fixed graph Laplacian.
- **Message Passing Neural Networks (MPNNs):** a general formulation (Gilmer et al. 2017) that covers many GNN variants. It defines message functions M and update functions U such that for each edge (i,j) a message $M(h_i, h_j, e_{ij})$ is computed (possibly using edge features e_{ij}), and then h_i is updated via $U(h_i, \sum_{j \in \mathcal{N}(i)} M_{ij})$. This abstraction is useful for implementing custom GNNs (like those with edge features or multi-step interactions).

Every GNN inherently assumes the graph structure is meaningful: by design, it **injects the relational inductive bias** that connected nodes should influence each other's state[6]. This is ideal for Gemantria's **concept graph**, where each node might represent a biblical concept or entity, and edges represent relationships (co-occurrence in verses, semantic similarity, etc.). A GNN can be used to compute **concept embeddings** – vector representations of each concept that encode the context of how that concept connects to others. These embeddings automatically reflect the graph's structure (e.g. two concepts with many shared neighbors will end up with similar embeddings). In practice, one could use a library like **PyTorch Geometric (PyG)** or **Deep Graph Library (DGL)** to implement this. *PyTorch Geometric* provides convenient GNN layers (e.g. GCNConv, GATConv) and handles graph mini-batches, making it easy to train GNNs in a few lines of code[8]. *DGL* similarly offers a flexible API (for PyTorch, TensorFlow, etc.) and optimized message-passing; it even includes specialized packages like **DGL-KE** for large-scale knowledge graph embeddings[9], which could be relevant if the concept graph is very large or has many relation types.

Using GNNs for Concept Graphs in Gemantria: By applying GNN models to the biblical concept graph, we can enhance relational reasoning in multiple ways:

- *Node Classification / Clustering*: If concepts are categorized (themes, people, places), a GNN can learn to classify each node's category by aggregating contextual clues from its neighbors. This yields **explainable results**: e.g. a concept is classified as "place" because it's connected to known locations, etc.
- *Link Prediction*: GNNs (or related graph embedding methods) can predict new edges or the strength of connections between concepts. This could reveal latent associations (e.g. two ideas that frequently appear together but weren't explicitly linked). Models like **Relational GCN (R-GCN)** extend GNNs to multi-relational graphs, learning separate weight transformations for each relation type[10]. This would allow Gemantria to handle different edge types (co-occurrence vs. topical similarity) when enriching the semantic network.
- *Concept Similarity and Analogy*: In the current pipeline, concept centrality and clustering are computed with graph algorithms. A learned GNN embedding could complement these by capturing higher-order patterns: e.g. two concepts that never co-occur but are connected via an intermediate concept might end up embedded closely. This provides a more nuanced measure of **graph-based similarity** than raw co-occurrence counts. Downstream, those embeddings can be used to find analogies or to feed into the search/rerank system (e.g. suggesting related concepts to explore).
- *Explainability via Attention*: By using an attention-based GNN (like GAT or Graph Attention with edge-type attention), we can inspect which neighbor concepts contributed most to a given concept's representation. This aligns with the goal of making the system's reasoning transparent – essentially highlighting the relational evidence for each inference.

In summary, geometric learning on graphs equips the Gemantria project with **trainable, structure-aware algorithms** that go beyond static network metrics. Instead of manually engineering features (like degrees or betweenness centrality), a GNN can *learn* features that are optimal for a task (forecasting, classification, etc.) while respecting the concept graph's connectivity. This is a powerful paradigm shift: the model can generalize and combine relational patterns in ways a fixed algorithm might miss, all while ensuring the outcomes are consistent with the known graph structure (permutation symmetry of nodes and the pattern of edges).

Learning on Manifolds and Non-Euclidean Spaces

While graphs handle discrete relations, **manifold-based geometric learning** deals with continuous geometric domains (curves, surfaces, or more abstract Riemannian spaces). Many data domains in science and engineering are manifolds – for example, 3D shapes (which are surfaces in \mathbb{R}^3), sensor networks on curved surfaces, or trajectories on a sphere. The challenge here is that **standard neural operations (like convolution)** are originally defined for Euclidean grids, and manifolds lack a global coordinate system or shift-invariance[7]. Geometric deep learning research has developed methods to generalize CNNs to manifolds and graphs, sometimes referred to as *non-Euclidean convolution*. Key approaches include:

- **Spectral Graph Convolution**: Interpreting a graph or manifold via the eigenfunctions of its Laplace operator. A convolution can be defined in the spectral domain by multiplying the signal's graph Fourier transform by a filter function. In practice, early methods (Bruna et al. 2013) used the graph Laplacian eigenbasis to perform convolutions on graphs and meshes. This provides a principled way to define frequency response on irregular domains, but it can be computationally heavy and non-local. Kipf & Welling's GCN can be seen as a

simplified spectral method where the filter is a linear function of the Laplacian (essentially a 1-hop averaging)[11].

- **Spatial (Patch-based) Convolution on Manifolds:** These methods define a local coordinate patch on the manifold and slide a kernel over it, analogous to image CNNs. Examples are **Geodesic CNN (Masci et al. 2015)** and **MeshCNN**, which pick a reference frame at a point (like using tangent plane or polar coordinates on the mesh) to apply a filter. The challenge is ensuring the filter orientation is consistent; this led to the concept of **gauge equivariance** (Cohen et al. 2019), where the CNN is invariant to how you choose local coordinates (akin to aligning the local chart before applying the kernel).
- **Hyperbolic and Spherical Embeddings:** In some cases, the “manifold” of interest is not a physical surface but a *geometry for embedding data*. For instance, hierarchical or tree-structured data can be embedded in **hyperbolic space**, which has properties suitable for representing trees with low distortion. Knowledge graphs or concept hierarchies could use *Poincaré embeddings* (Nickel & Kiela 2017) to capture the hierarchical relationships in fewer dimensions than Euclidean space. This is a geometric approach to representation learning: the model is constrained to place points on a curved manifold (like the hyperbolic disk), imparting an inductive bias that reflects hierarchy (common in ontologies or biblical genealogies).
- **Topological and algebraic techniques:** Beyond manifolds, techniques from algebraic topology (e.g. persistent homology, cell complexes) are being combined with neural networks to capture higher-order relationships (such as cycles or cavities in data). These are more advanced and beyond the standard scope of Bronstein’s 5G framework, but they exemplify the trend of incorporating more geometry to handle complex structures.

For the Gemantria project, direct applications of manifold learning might not be as obvious as graph learning, since our data (biblical text, concept networks, time series) is not inherently a curved geometric space. However, there are a few niche ways manifold-based thinking could help:

- **Temporal cycles as manifolds:** If temporal patterns have periodicity (e.g. thematic cycles or seasonal trends in mentions of concepts), one can model time modulo that period as a *circular manifold* (a 1D circle $\$S^1\$$). In practical terms, this means using features like $\$|\sin(2\pi t/T)|, |\cos(2\pi t/T)|\$$ for seasonality or designing a network that is equivariant to time shifts by $\$T\$$. This introduces a **geometric prior for time-series** – the network knows, for example, that a pattern repeating every 7 days is the same regardless of phase. Integrating such priors can improve forecasting of repeating patterns (perhaps relevant if Gemantria finds weekly liturgical cycles or other periodic trends in concept frequencies).
- **Knowledge graph embedding in curved space:** If the biblical concept graph has a hierarchy (for example, “Virtue” might branch into “Faith”, “Hope”, “Charity” etc.), embedding it in a non-Euclidean space like a hyperbolic manifold can preserve the hierarchical distances more faithfully than a flat space. Tools like the geoopt library or built-in functions in PyTorch Geometric allow optimization on Riemannian manifolds. This could be an advanced way to initialize concept embeddings such that general (broad) concepts are near the “center” of hyperbolic space and specific ones radiate outward, mirroring a semantic tree. GNNs or relational models could then refine these embeddings.

- **Manifold regularization for semantics:** When enriching biblical text semantically, we might treat semantic embedding spaces as manifolds. For example, word embeddings lie on a high-dimensional manifold. Techniques like aligning embeddings from different sources (biblical text vs. concept definitions) can be viewed through a geometric lens (e.g., using Procrustes alignment which is essentially finding a rotation in embedding space). Ensuring that a GNN or a transformer that links text to concepts does so smoothly (without tearing the embedding manifold) could improve consistency. This is more of a conceptual hint – practically, one can enforce that similar verses map to nearby concept embeddings on the manifold.

In summary, geometric deep learning on manifolds teaches us to respect the *geometry of data domains* beyond simple grids. Even if Gemantria doesn't deal with 3D shapes or physical manifolds, the mindset of **injecting geometric knowledge (like periodicity or hierarchy)** into models can inspire solutions. If a problem can be phrased as "data lives on a special space where usual neural nets don't directly apply," chances are a geometric DL approach or library exists to handle it. For instance, PyTorch Geometric extends beyond graphs to point clouds and meshes (point clouds can be seen as samples of a manifold)[8]. Should Gemantria ever need to model spatial relationships (maybe mapping biblical places geographically) or abstract semantic spaces, these tools would become relevant.

Group-Equivariant Neural Networks and Symmetry

A core theme of geometric AI is **leveraging symmetry**: if a task's outcome should not change under certain transformations of the input, we build that into the model. The classic example is the translational symmetry of images: a cat is a cat whether it's on the left or right of an image. A standard CNN is *translation-equivariant* – shifting the input causes an equivalent shift in output feature maps, thus the final classification is *translation-invariant*[4]. This idea generalizes to other groups of transformations.

Group-equivariant neural networks are designed so that if the input is transformed by some group element (rotation, reflection, permutation, etc.), the internal representation transforms in a predictable way (equivariance), or the output stays the same (invariance) for tasks like classification.

Examples of group symmetries and corresponding networks:

- **Rotational Equivariance (Images):** A network that processes images but is insensitive to rotation can be achieved by **Group CNNs**. Cohen & Welling's *Group-equivariant CNN* (2016) introduced convolution filters that are applied not just in shifted positions but also rotated orientations. Essentially, the filter bank is augmented to cover discrete rotations (e.g. 90° increments or even 8-fold rotations), and the convolution operation is defined over the group $\text{SE}(2)$ (position and orientation) instead of just the plane. The result is a feature map that responds similarly to a rotated version of the input, enabling the network to learn **rotation-invariant** detection without seeing all rotated examples in training. This is highly relevant to image or pattern recognition tasks where orientation is irrelevant.
- **Permutation Equivariance (Sets/Graphs):** As discussed, GNNs are inherently equivariant to node permutations – this is actually an instance of a group equivariant network, where the group is the symmetric group S_n acting on n nodes[12][5]. There is also the simpler case of DeepSets (Zaheer et al. 2017): a network that processes an input set $\{x_1, \dots, x_n\}$ by individually embedding elements and summing them is invariant to permutation. Graph

networks generalize this by also considering neighbors given an adjacency structure. In Gemantria's concept graphs, this equivariance is crucial: it means the analysis or predictions don't depend on arbitrary labeling of concepts, only on the relational structure.

- **SE(3)-Equivariant Networks (3D geometry):** For molecules or any 3D configuration, we often desire rotation and translation invariance (e.g. a molecule's properties don't depend on how you rotate the molecule in space). SE(3)-equivariant neural networks (such as *SE(3)-Transformer* or *E(3) equivariant GNNs* like **Tensor Field Networks** and **E3NN**) use spherical harmonics and careful tensor operations to ensure that when atomic coordinates are rotated, the learned representation rotates in the same way and the output (e.g. a scalar property like energy) is invariant[13]. These models have seen success in protein structure prediction and physics simulations, where respecting physical symmetries is essential for accuracy. A notable example is *AlphaFold2*'s attention architecture which incorporated rotational equivariance for handling protein 3D coordinates, a technique that contributed to its breakthrough in protein folding[13].
- **Other Symmetries:** If your data has other symmetry groups, similar principles apply. For example, **Graph isomorphism networks** tie into permutation symmetry in a different way (maximizing expressive power while remaining invariant). In text, one might consider **sequential equivariance** for time shifts or reversing if palindromic structure is irrelevant (though word order usually matters, so language models are not symmetric to permutation except in specific contexts like bag-of-words). In reinforcement learning on grids, you might want rotational or reflectional invariance of board states (applied in Go/Chess AI). In all cases, incorporating the symmetry can reduce the burden on learning and improve generalization.

Applying Group Equivariance in Gemantria: The biblical data might not obviously scream for group-theoretic symmetry as an image or physics task does, but there are opportunities to use these ideas:

- *Permutation-invariant set encoding:* When aggregating information from multiple sources (say, summarizing a set of verses relevant to a concept, or combining multiple concept embeddings to represent a cluster), using a permutation-invariant network (like a simple sum/mean of individual embeddings processed by an MLP) ensures the model's output doesn't depend on the order in which those inputs are provided. This is an easy win for explainability and correctness – e.g., a cluster's embedding should be the same regardless of the arbitrary ordering of concepts in that cluster.
- *Time-translation invariance:* For **temporal pattern analysis**, if we suspect that certain patterns repeat over time or that the exact start time of a pattern shouldn't matter, we could use architectures that emphasize **convolution in time** (temporal CNNs) or **Fourier features** that capture time-periodicity. A temporal convolution treats the time series with the same shift-invariance as an image CNN treats pixels – sliding a filter over time means the learned pattern can detect a subsequence anywhere in the timeline. This is effectively using the symmetry of the time axis (translation in time). If Gemantria is looking at narrative patterns or recurring themes across chapters, a temporal convolution or recurrence with shared weights already encodes the inductive bias that a motif is recognizable regardless of when it occurs. (Standard RNNs and transformers also share parameters across time positions, which is a form of equivariance to time-shift.)

- *Relational symmetries*: Some relations in the concept graph might be inherently symmetric or antisymmetric. For example, if an edge type means “A is associated with B,” that could be symmetric (A associated with B implies B with A), whereas a “is part of” relation would not be symmetric but might have other constraints (transitivity or such). While generic GNNs don’t enforce these logical constraints, one can inject such priors by *data augmentation* (adding reverse edges for symmetric relations) or specialized architectures. A **Relational GNN** could treat symmetric versus directed relations differently. Ensuring the GNN respects these properties (like by using an undirected edge for a symmetric relation) is akin to building domain symmetry into the model. This improves explainability because the model won’t arbitrarily break a symmetry that should hold in the knowledge graph.

Overall, group-equivariant neural networks offer a toolkit for tailoring models to *known invariances*, thereby reducing sample complexity and making learned representations more aligned with human expectations. In practice, using them in Gemantria might involve choosing the right architecture (e.g. using a circular padding or convolution for periodic time data, or using an attention mechanism that is invariant to input ordering for sets of verses). There are specialized libraries for equivariance: **e3nn** (in PyTorch) provides layers for 3D rotations; **LieConv** provides generic group convolution layers for certain continuous groups. However, many invariances can be handled with simpler means (data preprocessing or standard layers used cleverly). The main takeaway is to **identify any symmetry in your task and exploit it** – if your analysis shouldn’t change when two input elements swap or when an input is rotated, then you can likely simplify the model’s job by encoding that fact. This leads to models that are easier to train, more **data-efficient**, and often more **interpretable**, since their features align with symmetry modes (e.g. some features might respond to “pattern X regardless of phase” or “concept set Y regardless of order,” which matches how an expert might reason about the data).

Figure: Example of **geometric equivariance** in a CNN. In images (modeled on a 2D grid), a convolutional layer commutes with translations: shifting the input (by a translation $\$T\$$) yields a shifted feature map at the output. This *translation equivariance* is a built-in symmetry of CNNs – it means the model outputs the same detection no matter where an object is in the image[4]. By exploiting the grid’s shift symmetry, CNNs achieve locality and weight-sharing, making them far more efficient and generalizable for vision tasks than fully-connected networks.

Figure: Example of **geometric equivariance** in a GNN. In graph-structured data, the relevant symmetry is permutation of node indices (and potentially continuous transformations of node attributes). A message-passing network $\$F\$$ is designed such that permuting the input graph’s node order by a permutation matrix $\$P\$$ (red) and consistently permuting the adjacency $\$A\$$ to $\$PAP^T\$$ leaves the network output unchanged except for the same permutation applied to output node features[5]. If node features themselves have geometric meaning (e.g. positions that can rotate by $\$R\$$, shown in green), an equivariant GNN can be built to commute with those transformations as well. This means the GNN respects the graph’s structural and geometric symmetries: it doesn’t matter how nodes are labeled or oriented, the relational computations yield an appropriately transformed output. Such equivariant designs ensure that **only the relational structure and intrinsic patterns** affect the result, not arbitrary representations.

Relational Inductive Bias and Graph Networks

One of the motivations for geometric deep learning is to enable **relational reasoning** and **combinatorial generalization** in AI systems[14][6]. A classic deep learning model like a feed-forward network or vanilla transformer treats input data as a flat array – it has no explicit notion of entities and relations. In contrast, our *human* cognition system distinctly recognizes objects and the relations between them, allowing us to reason about novel combinations of known components. **Relational inductive bias** refers to designing models that **expect and leverage relationships** among entities. By structuring input as a graph (or more generally, a set of entities with connections) and using architectures that operate on graphs, we bias the model to develop functions that mirror relational logic (who is connected to whom, and how).

The work of Battaglia et al. (2018) introduced the concept of a **Graph Network** as a neural building block with a strong relational inductive bias[6]. A Graph Network processes a graph with *node attributes*, *edge attributes*, and even *global attributes*, updating all of them through learned functions. In each layer of a typical Graph Network:

1. **Edge update:** Each edge e_{ij} (from node i to j) updates its feature based on the previous features of the sender node i , receiver node j , and the edge itself (and possibly a global feature). This models pairwise interactions (e.g. how one concept influences another along a specific relation).
2. **Node update:** Each node then updates its feature based on its previous feature and the aggregated messages from incoming edges (after the edge updates). This is analogous to message-passing GNNs but with potentially richer edge computations from the prior step.
3. **Global update:** An optional global state can be updated from summary statistics of all nodes or edges (useful for graph-level tasks or context).

This framework generalizes many earlier graph models and is very expressive. Importantly, it separates the concern of *what is being computed* from *which entities are connected*, making it easy to encode known rules. For example, if you know a certain relation is transitive, you might design a multi-hop propagation or constrain the edge function accordingly. Battaglia's paper demonstrated that Graph Networks can solve complex reasoning tasks (like physics simulations, logical puzzles) in a way that extrapolates to larger systems than seen in training – a hallmark of combinatorial generalization[14][10]. By “**providing a straightforward interface for manipulating structured knowledge**”[10], graph networks serve as a bridge between symbolic reasoning and neural networks.

In Gemantria's context, adopting a relational inductive bias means explicitly treating biblical knowledge as a system of entities and relations (which is already being done with the concept graph) and using models that *compute over that system* rather than treating it as unstructured data. Some concrete recommendations:

- **Use Graph-structured representations everywhere feasible:** If there is a process currently done in an unstructured way (for instance, analyzing a list of verses or computing similarities by embedding text only), consider if it can be improved by incorporating the relational graph. For example, semantic enrichment of a verse could be done by linking it to concept nodes (forming a small bipartite graph of verse-to-concepts) and then using a graph network to propagate contextual signals: concepts that are related in the concept graph could reinforce each other's relevance to the verse. This would bias the system

to prefer interpretations consistent with known links (a form of structured disambiguation).

- **Graph Networks for Pattern Discovery:** The current pattern discovery in Gemantria involves clustering and centrality measures on the concept graph. A learned graph network could augment this by learning to identify subgraph patterns or communities that correlate with interesting outcomes (perhaps themes or emotional tones in the text). One could train a graph network to predict, say, which concept clusters are most “influential” or which subgraphs correspond to specific biblical themes, using some labeled data or even unsupervised objectives (like reconstructing connections). The relational inductive bias ensures the model looks at combinations of concepts rather than individual ones in isolation.
- **Combinatorial generalization in forecasting:** When forecasting temporal patterns, if we incorporate multiple series and their relations, we want the model to generalize to new sets of relations. A carefully designed relational model (like a graph neural network that can be applied to any graph of concepts) can learn a rule such as “if concept A strongly influences concept B and A spikes, then expect B to spike after a short delay.” At test time, the model could apply this rule even for a pair of concepts it never saw together in training, as long as the concept graph indicates a strong $A \rightarrow B$ relation. This is combinatorial generalization – using learned relational rules on new combinations of entities. Standard time-series models would struggle to do this, because they’d treat each pair or each multivariate series as distinct. A relational GNN approach shares the learned interaction dynamics across all pairs of connected nodes, thus **scaling insights to unseen scenarios**.
- **Explainability via structured reasoning:** With an explicit relational model, we can trace *why* a prediction was made in terms of which relationships were activated. For example, if a graph network predicts a future increase in concept “Faith” because “Love” had a surge earlier and the two are connected, we can inspect the messages passed along the $\text{Love} \rightarrow \text{Faith}$ edge. This is more interpretable than a black-box multivariate ARIMA, for instance, because it aligns with a causal storyline (“Love influences Faith”). The **relational graph serves as an explanation scaffold**.

In practice, implementing relational inductive bias in code can be done using the same GNN libraries (PyG or DGL). Both support building custom message-passing layers where you have access to edge attributes and global attributes. DGL’s `update_all` API or PyG’s `MessagePassing` class can be used to implement the three-step (edge, node, global) update as in Battaglia’s Graph Networks. There is also the DeepMind `GraphNets` library (for TensorFlow/Sonnet) that was released with the Battaglia paper, which provides high-level modules for graph networks[10]. However, PyTorch Geometric has since included similar capabilities (e.g., `MetaLayer` which allows you to define separate edge, node, global update functions).

To summarize, **relational inductive bias** is about **thinking in terms of entities and relations** and building models that naturally operate on graphs of those entities. By doing so in Gemantria, we ensure that as the scope of analysis grows (more concepts, new books, different combinations of themes), the system can handle it gracefully. It will treat a new relation just as it treated ones seen before – because it was designed to *expect relationships*. This not only improves generalization and reduces data requirements, but it also means the system’s behavior can be understood in relational

terms (which is how domain experts often conceptualize biblical themes: interconnected and context-dependent).

Tools and Libraries for Geometric Deep Learning

Implementing geometric deep learning techniques has been made much easier by various open-source libraries. Here we list some widely used tools and how they could support Gemantria's development:

- **PyTorch Geometric (PyG):** A popular library built on PyTorch for deep learning on irregular data like graphs, point clouds, and manifolds[8]. PyG provides efficient data structures for graphs and a large collection of GNN layers and utilities. For example, it offers ready-to-use implementations of GCN, GAT, GraphSAGE, global pooling layers, heterogeneous graph support, etc. PyG would allow rapid prototyping of a concept graph GNN – you can load the concept network into a Data object and apply standard layers to perform node classification or link prediction. Its integration with PyTorch means you can also incorporate text features (perhaps from a language model) and jointly train a model that combines text and graph inputs. PyG also has basic support for point clouds and meshes (though those might not be needed for Gemantria).
- **Deep Graph Library (DGL):** Another full-featured library for graph neural networks, which is framework-agnostic (works with PyTorch, TensorFlow, MXNet)[15]. DGL emphasizes performance and scalability – it's designed to handle graphs with millions of nodes/edges by efficient sampling and minibatching. One advantage for Gemantria is DGL's **heterogeneous graph** support: if your concept graph has different node types (e.g. concept nodes vs. verse/document nodes) and different edge types, DGL can handle that neatly. For example, you could create a bipartite graph of verses to concepts and run a heterogeneous GNN that updates verse representations based on concepts and vice versa. DGL also provides domain-specific extensions: **DGL-KE** for knowledge graph embeddings (TransE, DistMult, etc.) which could be used to pre-train embeddings of the concept graph; **DGL-LifeSci** for molecular graphs (with some equivariant models included). The choice between PyG and DGL often comes down to familiarity and specific needs – both are excellent. PyG might be slightly simpler for quick experiments in PyTorch, whereas DGL might shine if you have a very large graph or need multi-GPU training[16].
- **PyTorch Geometric Temporal:** An extension library specifically for **spatio-temporal graph learning**[17]. It includes implementations of many state-of-the-art models for dynamic graphs and time-series on graphs (e.g. DCRNN, STGCN, T-GCN, Graph WaveNet, etc.) out of the box. PyG Temporal provides data iterator classes for handling snapshots of temporal graph data and a variety of neural architectures that combine graph convolutions with RNNs or attention for forecasting tasks[17]. For Gemantria, this could be a gold mine for the **temporal pattern forecasting** part – instead of coding a custom GNN+LSTM, one could leverage these ready implementations. For example, **DCRNN (Diffusion Convolutional RNN)** which was designed for traffic forecasting could be repurposed: treat each concept as a “sensor” node and the concept graph as roads; the time series of concept mentions is analogous to traffic flow. DCRNN uses a diffusion graph convolution to mix information from neighbors and a sequence-to-sequence RNN to forecast future values[18]. This approach has been proven to handle both spatial and temporal dependencies effectively[18]. With PyG Temporal, one could set up a similar model in a few lines and see if concept co-occurrence links help predict surges or drops in related concepts.

PyG Temporal is also useful for *dynamic graphs*, where the edge structure might change over time (not sure if that's relevant for biblical data, but if you consider tracking which concepts are related at different points in the narrative, that could be dynamic).

- **Knowledge Graph and Heterogeneous Data Toolkits:** Beyond GNN-specific libraries, there are tools for dealing with knowledge graphs and multi-relational data. **PyKEEN** is a Python library for knowledge graph embedding algorithms (not GNN-based, but methods like TransE, RotatE, ComplEx which learn vector embeddings for entities and relations). If the concept graph is large and you want a quick way to get embeddings encoding each relation type's semantics, PyKEEN could be used to train an embedding model using a knowledge graph approach. However, increasingly GNN-based methods (like R-GCN) are competitive with or outperform static embeddings, especially when we want to integrate with other modalities. **HeteroData in PyG or heterograph in DGL** are specific data structures to manage multiple types of nodes/edges. Gemantria's semantic enrichment might involve heterogeneous graphs (e.g. scripture verses connected to concepts, concepts connected to other concepts, concepts connected to external resources like dictionaries). These libraries allow you to apply different GNN layers to different relation types or even entirely different architectures for each part of the graph, all within one unified model.
- **Libraries for Equivariance and Manifold learning:** If needed, specialized libraries like **e3nn** (for 3D equivariant networks) could be used for any task involving rotations (likely not needed for text, but perhaps if analyzing geometric patterns in text embeddings or so – probably overkill here). For manifold optimization (like hyperbolic embeddings), libraries such as **GeoOpt** or **Potato (Poincaré embeddings)** can be handy. These are more esoteric but available if the need arises.
- **NetworkX and graph analysis libraries:** It's worth mentioning that for non-learning tasks (just analyzing the graph structure), NetworkX is a pure Python library that Gemantria might already be using for computing centrality, clustering (the code snippets suggest usage of Louvain, centrality metrics, etc.). While NetworkX is not for deep learning, it can be used in tandem with GNNs – e.g., use NetworkX to compute some graph statistics as features, or to pre-process the graph (finding the largest components, etc.). The geometric learning doesn't replace these classical analyses; rather, it augments them. One can even combine them: e.g., a GNN could use node degree or centrality as part of its input features to inform the learning.

Below is a summary table of some **key libraries and their uses** for quick reference:

Library / Toolkit	Purpose	Notes for Use
PyTorch Geometric (PyG)	Graph neural network library for PyTorch.	Many pre-built GNN layers (GCN, GAT, etc.), handles graphs as input to PyTorch models[8]. Great for rapid development on small to medium graphs and integration with other PyTorch components (e.g. Transformers for text). Supports heterogeneous graphs via HeteroData.
Deep Graph Library (DGL)	Scalable GNN library (multi-framework).	Highly optimized for large graphs and minibatch training. Supports multi-relational graphs easily[9]. Offers DGL-KE for knowledge graph embeddings and other domain-specific modules. Good choice if

Library / Toolkit	Purpose	Notes for Use
PyG Temporal	Extension of PyG for spatio-temporal data.	concept graph is very large or one needs distributed training.
PyKEEN / DGL-KE	Knowledge graph embedding (non-GNN).	Implements popular graph-based forecasting and dynamic graph models[17]. Ideal for trying advanced temporal models (like DCRNN, STGCN) on concept time-series. Saves effort in coding custom sequences.
NetworkX / iGraph	Graph analysis (not learning).	Collection of algorithms for learning embeddings of entities/relations (TransE, RotatE, etc.). Could be used to initialize concept embeddings or as a baseline for link prediction. DGL-KE is optimized for huge KGs (millions of nodes/edges)[9].
e3nn (PyTorch)	Equivariant neural network toolkit (3D).	Useful for computing graph metrics (centrality, shortest paths, community detection). Can generate features or evaluate structural patterns. NetworkX is easy but slow for large graphs; igraph or Graph-tool are faster alternatives.
GeoOpt (Manifold Opt)	Riemannian optimization in PyTorch.	Provides building blocks for $E(3)$ -equivariant networks (useful if we had 3D structures, probably not needed for text). Mentioned for completeness regarding group equivariant libs.
		Enables placing tensors on manifolds (like hyperbolic space) and optimizing on those surfaces. Could be used if hyperbolic embeddings of the concept graph were desired.

Using these libraries, Gemantria's developers can implement geometric models efficiently. For instance, one could use PyG to prototype a GAT on the concept graph in a few hours, or use PyG Temporal's TemporalSignal loaders to feed concept time-series + graph structure into a pre-built forecasting model. The libraries handle the heavy lifting of backpropagation through combinatorial structures, so the team can focus on experimentation and interpretation rather than writing low-level graph code. Moreover, the existence of **pre-trained models and tutorials** in these ecosystems (e.g. PyG has examples for node classification on citation networks, DGL has a knowledge graph completion tutorial) can accelerate development by adapting proven recipes to the biblical data scenario.

Typical Architectures and Use Cases in Geometric Learning

It is useful to have a bird's-eye view of the common architectures in geometric deep learning, alongside their typical use cases and relevance to Gemantria. The table below organizes some of these architectures by domain:

Architecture / Model	Description	Example Use Cases	Relevance to Gemantria
GCN (Graph Conv Network)[11]	First-order spectral GNN; aggregates neighbor features with equal weights (after	Node classification, link prediction on homogeneous graphs (social networks,	Good baseline for concept graph tasks (e.g. classify concepts or predict

Architecture / Model	Description	Example Use Cases	Relevance to Gemantria
	normalization). Simple and fast.	citation networks).	new links). Can propagate information about concept importance through the graph.
GraphSAGE	Inductive GNN; learns how to sample and aggregate neighbors (mean, LSTM, etc.). Can generate embeddings for unseen nodes by aggregating their neighbors' features.	Large graphs where training is done on a sample and you need generalization to new nodes (e.g. Pinterest PinSAGE for recommendations).	Could be used if new concepts might be added on the fly, or to embed subgraphs of the Bible not seen in training (e.g. apocryphal texts) by their neighbor structure.
GAT (Graph Attention Net)	GNN with attention weights on edges[5]. Learns which neighbors are most relevant when updating a node's representation. Often improves accuracy on graphs with noisy or varying connectivity.	Citation networks, knowledge graphs (can attend over different relation importance), anything where some neighbors are more influential than others.	Highly relevant for explainability – e.g. highlight which related concept most influenced a prophecy concept's forecasting. Could improve concept embeddings by focusing on strong semantic links and down-weighting incidental co-occurrences.
R-GCN (Relational GCN)	Extension of GCN for multi-relational graphs; has separate weight matrices per edge type (with regularization to avoid blow-up of parameters). Allows learning from multi-relationship knowledge graphs.	Knowledge base completion (e.g. Freebase, WordNet link prediction), entity classification using rich relation types.	If the concept graph distinguishes different relation types (e.g. “is-a” vs “co-occurs-inverse” vs “thematic similarity”), R-GCN can utilize that information rather than blurring all connections together. Can be used to enrich concept representations with typed relational info.
Graph Transformers	Newer class of models applying transformer-style attention to	Learning on graphs where long-range interactions matter	Might be considered for concept graphs if

Architecture / Model	Description	Example Use Cases	Relevance to Gemantria
	graphs (treating all nodes as interacting, sometimes with masking by edges). These can capture long-range dependencies better than local GNNs, at cost of higher computation.	(e.g. quantum chemistry graphs, or program analysis where distant nodes correlate). Also suitable for fully connected graphs (Transformers are essentially fully-connected GATs with learned adjacency).	very global patterns need to be learned (every concept influencing every other). Possibly overkill, but could be useful if trying to mimic something like “wisdom of the crowd” interactions among many concepts at once. A simpler alternative is to increase GNN depth or use jumping knowledge networks to capture multi-hop relationships.
Spatio-Temporal GNN (e.g. DCRNN, STGCN)[18]	Combines graph convolution with temporal sequence modeling (RNN or temporal conv). DCRNN uses diffusion convolution + Seq2Seq; STGCN uses alternating temporal conv and graph conv layers. These models handle forecasting on graphs.	Traffic forecasting (sensor networks), environmental modeling (e.g. predicting weather at stations, where stations form a graph), any multivariate time series with an underlying graph structure.	Perfect fit for <i>temporal pattern forecasting</i> in Gemantria. By using concept co-occurrence graph as the connectivity, these models can predict how a surge in one concept might lead to ripples in connected concepts. Also useful for detecting anomalies: e.g. if a concept spiked without its neighbors also spiking, that might be a notable outlier event (suggesting an unusual context in the text).
Spherical CNN / Harmonic Networks	CNNs that operate on spherical data (like a globe) or use Fourier harmonics to handle rotations continuously. Achieve rotational	Recognizing patterns on spherical images (planetary data, omnidirectional images) or 3D model recognition by treating	Not directly applicable to text, but conceptually if we consider certain cyclic structures (like the “circle” of

Architecture / Model	Description	Example Use Cases	Relevance to Gemantria
	invariance on continuous groups by design.	them as spherical functions.	concepts or periodic time), analogous techniques could be applied (e.g. use Fourier features for periodicity). Mentioned here to represent manifold-specific architectures in GDL.
Equivariant Networks (SE(3), SO(3))	Networks like SE(3)-Transformer, LieConv, Tensor Field Networks that maintain equivariance under 3D rotations/reflections. Use tools like spherical harmonics, Bessel functions to encode geometric tensors.	Molecular property prediction (force fields, quantum energy), protein interface modeling, any task needing physical invariances.	Likely no direct use in Gemantria unless doing something like modeling a geometric embedding space. However, one notion is <i>permutation equivariance</i> , which GNNs already cover. If one imagines the concept graph embedding as living in a geometric space, any invariances there (maybe scale invariance of concept importance?) could be treated similarly.

This table is not exhaustive but covers the primary architectures that might come up. For Gemantria, the graph-based models (GCN, GAT, R-GCN) and the spatio-temporal models are the most pertinent. It's worth noting that many of these can be combined or extended: for example, one could build a **temporal GAT** (apply attention-based aggregation at each time step, with an RNN across time) if certain concepts influence others with varying strength over time. Indeed, one research direction is combining the strengths of transformers and GNNs to get the best of both worlds – e.g., using a transformer to propagate information globally on the graph in each time slice (so any concept can potentially attend to any other if needed) and then a recurrent structure over time. Some recent papers explore such hybrid models for time series on graphs[19].

Integrating Geometric Learning into Gemantria's Stack

Finally, let's focus explicitly on how to plug these geometric deep learning techniques into Gemantria's components: **(1) Concept Graph analysis, (2) Temporal Pattern Forecasting, and (3) Semantic Enrichment of biblical data.**

1. Enhancing Concept Graph Analysis with GNNs

Current state: Gemantria's concept graph is analyzed with tools like centrality measures and clustering (e.g. Louvain for communities, degree/eigenvector centrality for influence) as evidenced in the code. These provide useful static insights (most connected concepts, cluster of related ideas, etc.). However, they are hand-crafted and not directly tuned to any specific task or output – they serve as general analytics.

Geometric AI enhancement: Introduce **trainable graph models** on the concept network to learn embeddings or predictions that align with project goals (like identifying significant concept groupings, or predicting links that correspond to thematic connections). This doesn't mean discarding centrality metrics, but augmenting them with learned representations.

- **Concept Embeddings via GNN:** Use a GNN (GCN/GAT) to compute low-dimensional embeddings for each concept. These embeddings could be trained in an unsupervised manner, e.g., using a **reconstruction or contrastive objective**. One approach: try to reconstruct the adjacency neighborhood of each node (link prediction) – essentially training a GNN to predict which edges exist from a node's embedding (similar to node2vec/GraphSAGE unsupervised). Another approach: use random walks or graph similarity to generate pairs of nodes that should be close/distant and train the GNN with a contrastive loss (as in GraphSAGE unsupervised mode or Deep Graph Infomax). The result will be concept embeddings that capture multi-hop connectivity patterns (more comprehensive than raw centrality). These embeddings can then be clustered to find communities, or used to measure concept similarity for search/ranking. Because they are learned, they can encode more complex patterns – for instance, two nodes might never connect directly but belong to the same cluster because of a series of intermediate connections (the GNN can learn to place them nearby in embedding space if that helps reconstruction).
- **Link Prediction and New Relations:** With a trained GNN embedding, one can apply a simple link prediction by, say, dot product or an MLP on pair of node embeddings to score potential edges. This could help semantic enrichment by suggesting edges that were not explicitly in the data but make semantic sense (e.g. connecting two concepts that are used in similar contexts). For example, if “forgiveness” and “mercy” rarely co-occur explicitly but share many neighbors, a GNN might embed them such that the model scores a high affinity between them, flagging a possible relation. Domain experts could review such suggestions to expand the concept graph.
- **Node Classification or Attribute Prediction:** If concepts have attributes (like category, sentiment, or other metadata), a GNN can be trained to predict those, effectively propagating attribute information through the graph. In Bible analysis, maybe one could tag concepts as “people”, “places”, “virtues”, “doctrines”, etc., via a small labeled set, and use GNN to spread these labels to other nodes by learning from graph structure. This creates a more semantically **enriched graph** because unlabeled nodes get inferred labels based on neighbors. Even if explicit

classification isn't needed, the hidden layers of such a model often yield a good representation of concepts in terms of those categories.

- **Explainable reasoning on the graph:** If one incorporates an attention mechanism or analyzes message contributions, the GNN's decisions can be interpreted in terms of graph connections. For instance, GNNExplainer or integrated gradients could identify which neighbor nodes and edges had the most impact on a particular concept's embedding. If Gemantria builds an interface on top of this, a user could click on a concept and see highlighted related concepts that influenced its understanding. This aligns with a goal of explainability – it's like the model saying "Concept X is characterized strongly by its link to Concept Y and Z." If Y and Z have meanings, that provides an explanatory context for X.

Integration steps: The concept graph is likely stored in a database or as NetworkX object. To integrate GNN: 1. Prepare the graph data for PyG or DGL (edges list, initial features for nodes – which could be as simple as one-hot identities or some content-based feature if available, e.g. an embedding of the concept's name or description). 2. Choose a model – e.g. a 2-layer GCN or 2-layer GAT as a starting point. Define a training objective depending on available data: supervised (if any labels or known groupings exist, e.g. maybe we know some concepts belong to known categories) or unsupervised. 3. Train the model on the concept graph. This could be done offline and the embeddings cached (since the graph is relatively static for the Bible corpus). The training process can also incorporate the **relational inductive bias** by design (it inherently does, by using the graph). 4. Use the learned outputs in the pipeline: e.g. replace or complement the current clustering with clustering in embedding space, use distances in embedding space to supplement the "edge strength" metric, or feed these embeddings into the temporal model (see next section).

One interesting possibility: The concept graph GNN could be trained jointly with the temporal forecasting model (multi-task learning), so that the embeddings are optimized to not only reconstruct the graph but also to be useful for predicting time series correlations. This is advanced, but if successful, it means the concept representations learned would be especially good at explaining temporal patterns (for instance, grouping concepts that tend to rise and fall together).

2. Temporal Pattern Forecasting with Relational Graph Neural Networks

Current state: The forecasting in Gemantria is currently using simple statistical models (naïve last-value, simple moving average, ARIMA if available) on each time series independently[20][21]. This means each concept's frequency over time is forecasted without explicitly using information from other concepts. Some rudimentary global metrics (like average error, distribution of models used) are computed but no complex cross-series modeling.

Geometric AI enhancement: Leverage the concept graph to perform **spatio-temporal forecasting**, treating concept time series as interconnected. The premise is that concept occurrences are not independent – spikes or trends often involve groups of related concepts (for example, an increase in references to "famine" might correlate with increases in "pestilence" or "sword" if those themes co-occur in prophecy, to use a biblical motif). By using a graph-based forecasting model, we can capture such correlations and improve predictions, especially for concepts with sparse data where neighbors can act as signal boosters.

Concretely, **Graph-based time series models** like DCRNN, STGCN, or Graph Attention networks with temporal attention can be applied:

- **Diffusion Convolutional RNN (DCRNN):** As described, models traffic as a diffusion process on a road network[18]. For Gemantria, think of “information diffusion” on the concept network – if a certain concept becomes prevalent in a chapter, related concepts might be “diffused” into prominence in subsequent chapters or verses. DCRNN would handle this by using the adjacency matrix (or a weighted graph of concept similarities) in a diffusion convolution operation: effectively a weighted summation of neighbor values, analogous to a GCN, but designed to capture both incoming and outgoing influence (it uses bidirectional random walks)[22]. The RNN (GRU or similar) then handles the sequence aspect, maintaining a hidden state per concept (node) that evolves over time.
- **Spatio-Temporal GCN (STGCN):** This model (Yu et al. 2018) uses a mix of temporal convolution (TCN) layers and graph convolution layers. It might apply a temporal Conv1D (which looks at a window of past values for a concept) then a graph conv to spread that information to neighbors, and repeat. The result is a forecast for each node. This is a bit less dynamic than DCRNN (which can, in principle, propagate influences indefinitely through the RNN’s state), but it’s simpler to train (just feed-forward, no explicit RNN).
- **Temporal Graph Attention:** A variation could use attention over neighbors at the current time and possibly attention over time steps. There are models like ASTGCN (Attention based STGCN) that add learnable attention on both spatial and temporal dimensions. Such a model could, for example, learn that at a certain time lag, concept A strongly attends to concept B’s state one step before, indicating a one-step propagation delay.
- **Temporal Graph Networks (TGNs):** These are a bit different – they often refer to models for continuous-time dynamic graphs (with events), but some designs could apply. The idea is to update node representations whenever an event (e.g. co-occurrence of concepts at a time point) happens. If Gemantria needed to model something like evolving concept relationships over time, a TGN could be applicable (keeping a time-aware embedding per concept that evolves as new verses/events come in chronologically).

Implementation for Gemantria’s forecasting: 1. **Data preparation:** Represent the temporal data as a sequence of graph snapshots. For example, one could define the graph $\$G\$$ where nodes are concepts and weighted edges indicate concept-concept relatedness (from co-occurrence stats). Then for each time step (could be each chapter, or each book, or any consistent time index), attach the value (frequency count) of each concept at that time as the node feature $\$x_i(t)\$$. This yields a sequence $\{(G, X(t))\}_{t=1}^T$ where $\$X(t)\$$ is the matrix of node features (one feature: the count at time t)[23][24]. This can be fed into PyG Temporal iterators or manually processed. 2. **Choose a model architecture:** If using PyTorch Geometric Temporal, one can pick a provided model like DCRNN or TemporalConv or TGCN. Alternatively, one can build a custom model by combining a GNN layer and an RNN: e.g. at each time step, do $h(t) = \text{GNN}(X(t), A)$ to mix neighbor info, then feed $h(t)$ into an LSTM or GRU that carries it to the next step. The output can be through a linear layer to predict $\$X(t+1)\$$ or the next few values (for multi-step forecast). 3. **Training objective:** Typically, minimize the mean squared error (or MAE) between the predicted and actual concept counts over time. If forecasting multiple steps, use a sequence forecasting loss (summing MSE over each horizon step). One might also incorporate an auxiliary loss, like a regularization that the predictions respect some constraints (non-negativity, etc., though neural nets inherently can learn that for counts). 4. **Leverage exogenous inputs if available:** In some forecasting, external factors are used. For biblical text, an “external input” might be something like the progression of narrative or known events. Possibly we could input a

“time” embedding or an indicator of which book/chapter it is, to allow the model to learn, say, seasonal effects or book-specific dynamics. This is analogous to adding a calendar or location feature in traffic forecasting.

Benefits expected: By training a graph-aware model, patterns like “Concept B often rises shortly after Concept A in connected verses” can be captured. The model might learn a lag correlation mediated by the graph edge. Empirically, in traffic domains, graph-based models significantly outperform treating each sensor independently, especially when data is sparse or highly correlated[18]. We anticipate in biblical data, if a chapter heavily emphasizes one concept, the next chapter (or related cross-references) might emphasize related concepts – a relational model could predict such couplings, whereas independent ARIMA would miss them.

For example, consider the concepts “dream” and “vision” in a prophetic book: they might have alternating prominence – when a prophet *has a dream*, the text uses “dream”, but when *interpreting it*, it might refer to the *vision*. A relational model seeing a spike in “dream” might predict an upcoming spike in “vision”, given an edge linking the two (because historically they appear together in narratives). A univariate model might just see one spike and decay with no clue of the second concept.

Graph correlation & anomaly detection: Another advantage is the ability to detect **anomalies or pattern deviations**. If the model strongly expects two series to move together (due to graph link), but in observed data one spikes alone, that could flag an interesting event (maybe that occurrence of the concept is in an unusual context where the usual companion concept is absent). Gemantria could incorporate this by examining the residuals or attention weights. For instance, if concept A is normally correlated with B, but at time \$t\$ A surged and B did not (and the model erred there), that time point might be worth a closer look by researchers.

We recommend using **PyTorch Geometric Temporal** to rapidly experiment with these models. The library even includes example datasets and training scripts for traffic which can be adapted. If writing from scratch, ensure to use batching if needed (though if the dataset is just one long sequence, you might just train on that full sequence, perhaps using sliding windows).

3. Semantic Enrichment of Biblical Text via Geometric Techniques

Semantic enrichment involves linking raw biblical text to higher-level concepts and external knowledge, providing deeper interpretation aids. Currently, Gemantria uses concept extraction (possibly via NLP or manual tagging) and then enriches verses by attaching concepts. Geometric learning can contribute to this in a few ways:

- **Knowledge Graph Embeddings for Semantic Similarity:** By embedding the concept graph (as described earlier) and possibly integrating it with external knowledge (like a Bible ontology or WordNet-like semantic net), we can obtain vector representations that allow semantic queries. For example, given a new phrase or query, one could find which concept in the graph it’s closest to in embedding space, thereby enriching the query with related concepts. This is more robust than string matching. Foundational work like **TransE** (translational embeddings) treat relations in a knowledge graph as translations in vector space – such techniques could map biblical entities in a way that relational patterns (like *King* -> *Kingdom* similar to *Prophet* -> *Prophecy*) are encoded as vector differences. A geometric perspective here is that these embeddings often lie in a space (potentially non-Euclidean) that reflects the structure of the knowledge graph.

- **Heterogeneous Graph Neural Networks combining text and concepts:** Imagine a graph where we have two types of nodes: **Verse nodes** and **Concept nodes**. An edge connects a verse to a concept if that concept is present in the verse. This is a typical bipartite graph representation of a corpus with an ontology. Running a **heterogeneous GNN** on this structure can enrich both sides: verses' embeddings get influenced by the concepts (and through them, by other related verses), and concept embeddings get influenced by the contexts of verses. This is akin to what papers on graph-based semi-supervised learning for text do: using a graph of documents and terms to propagate labels or topics. In Gemantria, one could use a label like "chapter theme" or "topic cluster" as a supervision on verse nodes and let the GNN propagate that information to concept nodes (so you learn concept topical categories), or vice versa if concept categories are known and you want to classify verses. Even without explicit labels, one can do **graph-based smoothing of embeddings**: for instance, start with an embedding for each verse (say from a language model or bag-of-words) and an embedding for each concept (maybe random or based on a definition). Then apply a couple of rounds of message passing: verse -> concept -> verse, etc. Concepts will aggregate the contexts of verses they appear in, and verses will in turn get an enriched representation influenced by related verses (two verses connected via a concept will exchange information). This can yield a kind of **semantic enrichment via graph context**. It's similar to a recommendation system idea (concepts are like "items" and verses are like "users" who "consume" items; a GNN can learn embeddings such that if two verses share many concepts, they end up similar, etc.).
- **Relational Reasoning for Interpretation:** Sometimes biblical interpretation involves **following chains of reasoning**: concept A in verse 1 relates to concept B in verse 2, which fulfills concept C in verse 3, and so on. A graph neural network that is run iteratively can naturally perform multi-hop reasoning. Each message-passing step is like following a link. If we had a task like question answering or cross-referencing (e.g. find verses related to this concept chain), a GNN could be set up to do a few steps of inference. For example, given an initial activation on a node (representing a query concept), the GNN could propagate that activation outward and highlight nodes (concepts or verses) two or three hops away that get a high score, effectively enumerating a relational inference. While Gemantria might not need a neural net to do simple graph traversal, the advantage is if the relations have weights or the reasoning needs to combine textual similarity and graph links, a learned approach can balance these signals optimally. It also can incorporate **logical rules as features** (some works use differentiable logic or constraints on GNN outputs to enforce, say, that certain paths correspond to known analogies).
- **Explainable Semantic Linking:** Using attention or specific architectures, we can make the semantic enrichment explainable. For instance, a verse embedding can be seen as a weighted sum of concept embeddings that occur in it (that's essentially a one-layer GNN from concepts to verse). Those weights tell you which concepts were most representative of that verse's meaning. If you further allow concepts to gather info from other concepts (a second layer), then a concept not explicitly in the verse could still influence if it's connected to one that is – which might surface implicit themes. But through attention weights you could trace: Verse -> Concept1 (explicit) -> Concept2 (connected, implicit). Such a chain could be offered as an explanation: "Verse talks about *Concept1*, which is closely related to *Concept2*, suggesting an implicit theme of *Concept2*."

Implementation ideas for semantic enrichment: - Construct the bipartite graph of verses and concepts. Use a two-layer heterogeneous GNN: first layer: concept-to-verse message (aggregating all concept vectors in a verse to update verse vector), second layer: verse-to-concept (aggregating all verses that mention a concept to update the concept vector). Use residual connections if needed. This will produce refined embeddings. If any labeled data exists (like tagged topics for verses, or known categories for concepts), you can train this in a semi-supervised way (minimize classification loss on labeled nodes). If no labels, you might train it as an autoencoder: try to reconstruct links or content (e.g. predict bag-of-words of a verse from the concept-enhanced verse embedding). - Use the output embeddings in search or analysis. For example, when a user searches a term, instead of matching verses by raw text, map the term to the nearest concept or concept embedding, then find verses whose enriched embeddings are similar – this might retrieve verses that don't contain the exact term but are conceptually related (a classic semantic search scenario). - The graph approach here injects **external knowledge (the concept graph)** into the purely text-based domain. This is known to improve results in many NLP tasks, because the model can overcome data sparsity and resolve ambiguities using the graph's information. For example, the word "John" could refer to a person or a book; connecting the verse to the concept "John (Person)" vs "Gospel of John" via context and the graph could disambiguate which concept node it links to, leading to more accurate indexing.

- As a concrete tool: DGL is quite adept at heterogeneous graphs. One could define a heterograph with two node types and one edge type ("verse–contains–concept"). DGL's built-in hetero GNN layers (like `dgl.nn.HeteroGraphConv`) allow you to specify a GNN on each relation. In this simple case, it's one relation, so it's straightforward. PyG also has `HeteroConv` or one can treat the bipartite graph as a normal graph by converting concept nodes and verse nodes into one big graph (just keep in mind which is which if needed).

Potential outcome: Semantic enrichment powered by geometric learning would mean that the **concept graph and the text inform each other**. Concepts get richer definitions (in vector form) based on actual usage in verses, and verses get deeper tagging based on the concept network. This two-way enrichment can surface themes and connections that might be missed if treating text or graph alone. It also naturally handles multi-lingual or abstract queries if the concept graph contains those links (because you're not matching literal text, but concepts). The approach remains **explainable** because the graph provides a rationale (we can always trace which concepts linked a verse to a query).

Using Geometric Priors in Pattern Detection and Analysis

"Geometric priors" refer to assumptions like smoothness, symmetry, or locality that we expect our data or desired solution to have. Employing these in *pattern detection* and *correlation analysis* will refine Gemantria's analytics:

- **Graph smoothness for correlation:** A common prior on graphs is that signals vary *smoothly* over connected nodes – i.e., adjacent nodes in the graph are likely to have correlated values (this is related to low-frequency components on the graph Fourier basis). In pattern detection, if we assume concept frequencies are correlated along the concept graph, we can apply techniques like **graph Laplacian smoothing** to denoise the time series or detect anomalies. For example, one could compute at each time t a "smoothed" concept frequency vector $\mathbf{X}_{\text{smooth}}(t) = (\mathbf{I} - \alpha \mathbf{L})^{-1} \mathbf{X}(t)$ where \mathbf{L} is the graph Laplacian and α a small weight. This essentially spreads each concept's value to its

neighbors. Patterns (like peaks) that are genuine and supported by related concepts will persist under smoothing, whereas isolated spikes (perhaps noise or transcription errors in data) might dissipate. This prior could be integrated as a layer in a neural model or used as a preprocessing for pattern detection.

- **Geometric anomaly detection:** Conversely, looking at **high-frequency components** (signals that change sharply across edges) can identify interesting patterns like concept outliers or surprising contextual combinations. A GNN can be trained (unsupervised) to flag anomalies by reconstructing expected values and seeing where errors are large. If we incorporate the graph prior, the model's "expected" pattern for a concept might be partly based on neighbor concepts. Large deviation means something new happened (e.g. a concept appears without its usual context).
- **Equivariance in pattern mining:** If searching for subgraph patterns (motifs) in the concept network, we can use the idea of graph isomorphism invariance: a pattern should be recognized regardless of where it occurs. Techniques like **Graph neural network motif detectors** (using GNNs to represent subgraphs and cluster them) inherently treat isomorphic subgraphs as the same (due to the GNN's permutation invariance). This could help in finding recurring relational patterns, like "Concept A, B, C form a triangle of mutual references" which might indicate a doctrinal trio or a narrative trope. A GNN-based approach can learn which subgraph structures correlate with something (e.g. maybe triangles vs chains might correlate with different semantic outcomes).
- **Cyclic or rotational priors in sequences:** If certain patterns in time are cyclic (7-year cycles, etc., which do appear in biblical narratives), one can incorporate a prior by using a *circular convolution* or a cyclic embedding of time (as discussed earlier with manifolds). For example, one could add features to the model like "year mod 7" or even create a second time axis that is mod 7 and convolve along that. This ensures the model notices patterns that repeat every 7 units even if the phase is different. That's a form of equivariance to a time shift of 7. If Gemantria is analyzing something like the pattern of sabbatical years or jubilee cycles, this is directly applicable.
- **Attention to geometry in evaluation:** When evaluating pattern forecasts or cluster quality, geometric considerations can help. For instance, measuring how well a clustering conforms to the graph connectivity (e.g. using modularity, which is a geometric measure on the graph) can tell if the learned representations are aligning with known structure. If not, one might adjust the model to pay more attention to graph links (stronger regularization towards smoothness, or adding a loss term that penalizes very different embeddings on connected nodes).

In practice, using these priors might involve adding regularization terms to a loss function (like a Laplacian regularizer $\lambda \sum_{(i,j) \in E} |h_i - h_j|^2$ to enforce smooth embeddings), or designing the architecture (like using convolution and pooling on graphs to detect motifs of certain shapes). PyG and DGL allow relatively easy implementation of such regularizers since you can iterate edges or use built-in Laplacian functions.

Concrete example: Suppose we want to improve pattern discovery of concept clusters. We could take the concept embeddings from a GNN and run a community detection algorithm on them (like k -means in embedding space). To encourage these clusters to align with graph communities, we could train the GNN embedding with a term that encourages that outcome – e.g., the **DeepGraphInfomax** approach uses mutual information to ensure the representation captures global graph structure. Another

approach is to use a **contrastive loss** where positive pairs are connected concepts and negative pairs are far apart in the graph. This way, the GNN is directly optimizing a geometric prior: nearby in graph -> nearby in embedding. The resulting clusters in embedding should coincide more with actual connected communities (likely meaningful themes).

In summary, geometric priors act like a compass, keeping the learning oriented toward **meaningful structure** rather than spurious patterns. By integrating such priors, whether through architecture (equivariant layers, convolution on graphs) or through loss functions (smoothness, invariance constraints), Gemantria's system will become more robust in detecting true patterns and ignoring noise. The outputs will also be **more interpretable**, since they align with human-understandable structures (graph communities, periodic cycles, symmetric relationships) rather than arbitrary parameters.

Conclusion

Geometric deep learning – “the era of Geometric AI” as Bronstein calls it – offers a powerful set of tools and principles for building AI systems that are **structure-aware, relational, and explainable**. By embracing these techniques, the Gemantria project can transform its analysis of biblical texts:

- **Graph Neural Networks** will imbue the concept network with learning capabilities, allowing the system to discover nuanced relationships and reasoning paths in the data, while maintaining transparency through the graph structure[10].
- **Manifold and symmetry-based methods** ensure that any known invariances (such as periodicities or known equivalences) are respected by the model, reducing noise and focusing the pattern discovery on meaningful variations[25].
- **Relational inductive biases** mean the models inherently treat “the Bible as a network of interconnected ideas” – which aligns perfectly with how scholars study scriptural themes – thus enabling combinatorial generalization such as applying learned patterns from one set of concepts to another[6].

Concretely, integrating geometric learning into Gemantria’s stack would proceed as follows: 1. **Concept Graph GNN module**: producing enriched concept representations and new insights (like predicted links or communities) from the raw graph[5]. 2. **Spatio-Temporal forecasting module**: combining the concept graph with time-series data to forecast and detect patterns that involve multiple concepts in tandem[18]. 3.

Heterogeneous graph semantic module: linking scripture text and concept knowledge in a joint model, to mutually enhance understanding of both and allow explainable semantic queries.

Throughout these, we maintain a focus on **explainability**: each recommendation (attention weights, graph paths, invariant features) can be traced back to an interpretable element (a relation in the graph, a symmetry in time, a cluster of concepts), addressing the need for transparency in theological and literary analysis.

By grounding the AI’s reasoning in the “geometric” structure of the biblical data (the graphs of concepts and relations, the sequences of narrative, the known symmetries of themes), Gemantria can achieve deeper insights while still ensuring that the results are **human-comprehensible and verifiable** – ultimately using state-of-the-art geometric deep learning to augment age-old methods of biblical scholarship with modern machine intelligence.

Sources:

- Bronstein et al., *Geometric Deep Learning: Going beyond Euclidean data*, 2017 – provided the foundational survey of GDL on graphs and manifolds[26].
- Bronstein et al., *Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges*, 2021 – unified mathematical framework for modern architectures, emphasizing symmetry (Felix Klein's Erlangen Programme)[2][12].
- Battaglia et al., *Relational inductive biases, deep learning, and graph networks*, 2018 – introduced graph networks and argued for building AI with structured representations for combinatorial generalization[6].
- Kipf & Welling, *Semi-Supervised Classification with Graph Convolutional Networks*, ICLR 2017 – introduced the GCN model widely used for graph node classification.
- Veličković et al., *Graph Attention Networks*, ICLR 2018 – introduced attention mechanisms in GNNs for better weighting of neighbors (useful for explainability).
- Schlichtkrull et al., *Modeling Relational Data with Graph Convolutional Networks*, ESWC 2018 – introduced R-GCN for multi-relational knowledge graphs.
- Li et al., *Diffusion Convolutional Recurrent Neural Network for Traffic Forecasting*, ICLR 2018 – an application of spatio-temporal GNN for time series, relevant to concept forecasting[18].
- Yu et al., *Spatio-Temporal Graph Convolutional Networks*, 2018 – another approach to graph-based time series forecasting using purely convolutions.
- Various tool documentation: PyTorch Geometric[8], DGL[9], PyG Temporal[17], which provide practical implementations of the above techniques.

By following the guidance and examples above, the Gemantria team can integrate these geometric deep learning components step by step – first at the data preprocessing and embedding level, then into predictive modeling – and iteratively improve the system’s ability to recognize and reason about the rich tapestry of connections in biblical literature. The result will be a more **structured AI**, one that “understands” the Bible not as isolated verses or words, but as a connected graph of concepts and patterns, much like a scholar might – fulfilling the promise of geometric AI in a very meaningful domain.

[3][6]

[1] [3] [4] [5] [12] [13] [25] Towards Geometric Deep Learning

<https://thegradient.pub/towards-geometric-deep-learning/>

[2] [2104.13478] Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges

<https://arxiv.org/abs/2104.13478>

[6] [10] [14] [1806.01261] Relational inductive biases, deep learning, and graph networks

<https://arxiv.org/abs/1806.01261>

[7] [11] [26] graphics.stanford.edu

http://graphics.stanford.edu/courses/cs233-25-spring/ReferencedPapers/GCNN_Geometric%20deep%20learning-%20going%20beyond%20Euclidean%20data.pdf

[8] PyTorch

<https://pytorch.org/>

[9] [15] [16] Deep Graph Library

<https://www.dgl.ai/>

[17] PyTorch Geometric Temporal Documentation — PyTorch Geometric Temporal documentation

<https://pytorch-geometric-temporal.readthedocs.io/en/latest/>

[18] [22] [1707.01926] Diffusion Convolutional Recurrent Neural Network: Data-Driven Traffic Forecasting

<https://arxiv.org/abs/1707.01926>

[19] Temporal Edge Regression with PyTorch Geometric - Towards AI

<https://pub.towardsai.net/temporal-edge-regression-with-pytorch-geometric-3267d79cfa03>

[20] [21] export_stats.py

<file:///file-E5zs1wzMMbBuLZvg3uKRU>

[23] [24] PyTorch Geometric Temporal: Spatiotemporal Signal Processing with Neural Machine Learning Models

<https://arxiv.org/pdf/2104.07788.pdf>