



# Analysis of Repeated Work and Branch Management in Gemantria

## Branch Redundancies and Repeated Efforts

**Branch Proliferation:** The Gemantria repository shows an **overabundance of feature and fix branches**, many of which indicate duplicate or iterative work on the same tasks. A recent branch health report lists **over 150 branches** (both local and origin) active or stale. For example, the “**ops-stabilize-package-guards-017” feature was implemented multiple times**, evidenced by branches named `ops-fast-software-integrity-017a`, `ops-stabilize-package-guards-017`, and even a “clean” redo branch `ops-stabilize-package-guards-017-clean` <sup>1</sup>. This suggests that feature #17 had to be reworked more than once, likely due to complications or mistakes in earlier attempts.

**Multiple Attempts per Task:** Similarly, other tasks have letter-suffixed branch names (e.g. `024b`, `024c` for a “**make duplicate-target guard**” fix), indicating **partial fixes and follow-up fixes**. In one case, a CI branch `ci/make-dupe-guard-024b` was followed by another branch `024c-make-dupe-guard-main`, implying the initial fix wasn’t sufficient and had to be redone <sup>2</sup> <sup>3</sup>. These repeated attempts show that changes are being merged without being fully correct, necessitating additional “b” or “clean” branches to tidy up or complete the work.

**Integration Branch Duplication:** The repository also maintained **“integration” branches that duplicate feature branches**, which is an unusual pattern. For instance, after finishing feature branch `feature/pr-002-bible-ro`, a corresponding `integrate/feature/pr-002-bible-ro` branch was created <sup>4</sup> <sup>5</sup>. This happened for multiple features (e.g. PR-003, PR-004, etc.), effectively doubling the branch count for the same piece of work. Using separate integration branches to merge feature branches adds overhead and can lead to **duplicate merges or conflicts** – in one commit, the team had to “*dedupe*” changes by resetting files to main’s state while integrating `infra-guards-001` <sup>5</sup>. This duplication of effort suggests a pain point in the Git workflow.

**Phase Branch Overload:** Development was organized into phases, but the approach became fragmented. **Phase 8** work, for example, was split into numerous micro-branches (`phase8-kickoff-001` through `phase8-closeout-018`, among many others) for individual sub-tasks like anomalies logs, badge embedding, integrity checks, etc. <sup>6</sup> <sup>7</sup>. Many of these were marked “*(local-only)*”, meaning they were not independently merged back – instead, they were likely combined in a final “closeout” branch. This granular branching caused **integration challenges**, as evidenced by the need for a final Phase-8 closeout merge <sup>8</sup>. Similarly, **Phase 9** had a bundle of branches for various sub-tasks (centrality, audit, quality gates), even showing two nearly identical branches (`phase9-edge-audit-anomalies-badge-026` vs. `phase9-edge-audit-anomaly-badges-026`) – a sign that one branch superseded another, implying rework due to perhaps a naming error or approach change. By Phase 14, the project even needed a special “**phase14-rescue**” branch to salvage that phase’s objectives, and only then a “Phase 14 Complete” branch

was finished – a clear indication that **the first attempt at Phase 14 went awry and had to be rescued and redone**.

## Underlying Pain Points

Several root causes emerge for the repeated work and branch redundancy:

- **Overly Complex Git Workflow:** The practice of creating separate *integration branches* for each feature and maintaining so many concurrent branches led to confusion and duplicate work. Merging feature branches into intermediate integration branches (instead of directly into main) increases the chance of conflicts and misalignment. This complexity made the team effectively redo merges and conflict resolution multiple times, rather than once.
- **Lack of Early Consolidation:** Because features were developed in isolation (many local “Phase” sub-branches) and only consolidated at the end of a phase, integration problems were discovered late. This “*big bang*” *integration* style resulted in emergency fixes and duplicate efforts. For example, multiple Phase 8 branches had to be reconciled in a final step, and Phase 14 required a late “**rescue effort**”, indicating that issues were found only at phase-end rather than earlier in the process.
- **Quality Gates Causing Rework:** The project has very strict CI and governance rules (e.g. code coverage  $\geq 98\%$ , linting, determinism checks). While high standards are good, in practice the team sometimes bypassed or deferred them, causing later rework. For instance, at one point a branch temporarily **disabled lint and checks to “unblock” a merge** <sup>9</sup>. This allowed code with issues into main, but then multiple follow-up “fix/ci-linting-wave” branches were needed to clean up technical debt (reducing hundreds of linter errors in waves) <sup>10</sup>. In short, skipping quality steps early led to significant redoing of work later to get back up to standard.
- **Repeated Decision Reversals:** Some rework came from **inconsistent decisions or requirements churn**. A striking example is the **Ketiv vs. Qere gematria policy**. The project decided in *Phase 2* (back in early 2025) via ADR-002 that gematria calculations use **Ketiv (written form) as primary**, rejecting the alternative Qere-first approach <sup>11</sup> <sup>12</sup>. However, during *Phase 13* a proposal resurfaced to prioritize Qere, which had to be explicitly rejected as it conflicted with the established policy <sup>12</sup>. This indicates a **knowledge management lapse** – the earlier decision wasn’t communicated or remembered, causing the team to revisit and debate a solved issue. Such reversals waste time and effort, effectively “redoing” analytical work.
- **Unfinished or Sloppy Initial Implementations:** The presence of “-clean” redo branches and letter-suffixed attempts suggests *initial implementations were sometimes rushed or not fully thought through*. For example, feature 017 required a **clean refactor branch\*** to stabilize the pipeline <sup>1</sup>, implying the first pass was too messy to fix incrementally. This reflects a pain point in execution – perhaps pressure to deliver quickly led to technical debt that then required a fresh start to sort out.
- **Branch Management and Stale Work:** The branch health report shows many **stale local branches** (local Y / remote N), meaning branches that were not properly closed or deleted after merge <sup>13</sup>. This clutter makes it hard to know which work is truly done and which is pending, increasing the risk of duplicating efforts because developers might spin up a new branch not realizing a similar attempt

exists. The fact that a dedicated script `branch_health_report.py` was written to assess branch status <sup>14</sup> <sup>15</sup> highlights this as a known pain point – the team is spending effort just to track the tangle of branches.

## Strategies to Avoid Repeating Work

To recover from the current state and prevent these issues going forward, several changes are advisable:

- **Simplify the Git Workflow:** Move towards a more straightforward branching model. For example, consider adopting a **trunk-based development** or a simpler feature branch model where **feature branches are merged directly into `main` (or a dev trunk) once completed**, then deleted. Eliminate the extra “integrate/feature/\*” branches – they are unnecessary indirection. Instead, use pull requests to merge each feature branch to the mainline and run CI on each PR. This way, integration happens continuously, and problems surface early. The team will avoid duplicating merge efforts and cut down on stale branches, since each feature branch has one life cycle (open PR → review → merge → delete).
- **Integrate Early and Often:** Break the habit of long phase-ending integrations. Instead of accumulating dozens of micro-branches for a phase and merging at the end, aim to **merge incremental progress regularly**. If Phase 8 required 18 sub-tasks, the team could have merged intermediate results as soon as each was coded and tested, behind feature flags if necessary. Frequent integration ensures that conflicts and system-wide issues are discovered and resolved in small pieces, rather than in a large “rescue” mission. This approach would likely have prevented a scenario like Phase 14’s rescue, because issues would have been caught before the entire phase went off-track.
- **Strengthen Planning and Knowledge Sharing:** The project should enforce using the **single source of truth (SSOT) documentation and ADRs** for all significant decisions. Before implementing a new idea, developers and the PM should verify if it was considered previously. For instance, documenting the Ketiv/Qere decision in the SSOT and referencing it in planning docs helped clarify the conflict <sup>12</sup>; going forward, ensure every team member is aware of these records. Regular briefings or a living **Master Plan** (which Gemantria has) should be consulted so that phases don’t re-litigate settled matters. This will stop the team from unknowingly redoing conceptual work and prevent contradictory implementations.
- **Improve Initial Quality to Reduce Rework:** It’s critical to “**build it right the first time**” on each feature to avoid a cascade of fix branches. The PM (myself) should adjust timelines or expectations so developers can write tests, run linters, and ensure a feature meets standards *before* merging. We saw that bypassing linting or tests to rush a merge led to waves of cleanup <sup>10</sup>. Instead, we should treat the CI gates not as obstacles to work around, but as part of the definition of “done” for each task. If a PR is failing checks, keep it open and fix the issues instead of merging and fixing later. This discipline will drastically cut down the “redo” branches. It might slow down individual merges slightly, but saves time overall by **avoiding rework**.
- **Incremental Quality Improvement:** On the flip side, if the codebase already has quality issues (e.g. many lint errors or low coverage), create a prioritized plan to fix these **in a controlled, incremental**

**way** rather than ad-hoc. For example, schedule dedicated refactoring days or specific “quality improvement” tasks, and track them. This way, improvements like the lint fixes (which were done in multiple waves) are planned and reviewed, instead of emergency patches. A controlled approach prevents chaotic multiple branches addressing the same problem.

- **Clean Up and Refocus:** As an immediate recovery step, **prune the unused branches** to tidy the repository. Many branches listed as only local or already merged can be deleted to reduce confusion. Next, audit any pairs of branches where work was duplicated (for instance, ensure only the final “-clean” or “-rescue” branch remains relevant and the older attempt is fully superseded). This cleanup will help developers focus on the current code and reduce the cognitive load of the repository’s history.
- **Limit Parallel Work in One Area:** The explosion of Phase 8 and Phase 9 sub-branches hints that multiple team members or agents tackled a bunch of related tasks in parallel. While parallelism can speed up progress, it needs coordination. Introduce better **coordination or locking of related tasks** – for example, use a single branch for closely related changes or at least ensure integration after a few related branches, before starting more. A “**feature lead**” could oversee each phase to prevent duplicate or conflicting efforts. By sequencing work more and avoiding dozens of simultaneous tiny branches, we won’t end up redoing work to reconcile them later.
- **Role of PM and Communication:** As the PM, I acknowledge that rapid phase-driven development and aggressive timelines have contributed to these pain points. Going forward, I will prioritize **clear communication of priorities and changes** to avoid mid-phase pivots that cause thrash. We will hold brief **post-mortems at the end of each phase** to identify what went wrong (for example, why Phase 14 failed initially) and ensure we apply those lessons immediately. I will also adjust the planning: if we notice repeated attempts (like the “017” feature saga), that’s a signal to pause and address root causes (was the spec unclear? Was the scope too large? Do we need expert help?) before simply trying again. This more thoughtful management approach will help break the cycle of doing the same work multiple times.

## Recovery Plan Moving Forward

**Immediate Next Steps:** First, we should merge or close all redundant branches. For example, finalize the `ops-stabilize-package-guards-017-clean` branch (which contains the polished fix) and ensure it’s deployed, then **delete the older 017 and 017a branches** to avoid any confusion <sup>1</sup>. Perform a similar consolidation for any “a/b” branches like the `024c` duplicate-target guard – only the final solution should remain active <sup>2</sup> <sup>3</sup>. It’s also wise to run the branch health script again and formally archive or remove branches marked as local-only experiments from Phase 8/9 that are already incorporated in main. This cleanup will give us a clean slate.

**Refocus on Core Goals:** Next, revisit the **MASTER\_PLAN** and current phase goals with the team. We need to clearly identify any remaining blockers or P0 tasks and tackle them one by one with the improved workflow. Any ongoing “rescue” efforts (e.g. if Phase 15 is in progress) should be stabilized by applying the lessons above (small merges, no new parallel branches until the rescue is merged). The key is to **finish what’s started before moving on**, to restore confidence that our main branch is up-to-date and correct.

**Process Changes:** We will implement a streamlined Git policy as discussed – no more integration branches, and stricter branch naming conventions to avoid confusion. Perhaps adopt a format that ties branches to issue IDs or user stories rather than these phase numbers that got reused. This can prevent mix-ups like two branches both numbered “026” for different purposes. We’ll document this in an updated CONTRIBUTING guide so everyone follows the same process.

**Continuous Integration and Review:** Going forward, every pull request must pass CI checks **before** merge. If a check is overly noisy or blocking progress (as happened with lint), we will address it by configuring the tool (e.g. adjusting lint rules or using incremental linting) rather than bypassing it. The PM (myself) and tech leads will monitor PRs for any sign of “quick fix now, quality later” and push back on merging such PRs. It’s cheaper to correct issues in the original branch than to spin up a fix later.

**Knowledge Management:** We’ll also make better use of the documentation already in place. The **SSOT/MASTER\_PLAN.md** and ADRs should be kept up-to-date and **consulted at each phase planning** meeting [16](#) [12](#). If a new idea comes up, we’ll double-check prior phases’ decisions to ensure we’re not covering old ground. Encouraging a culture of reading the docs (perhaps assigning someone to be an “ADR champion” to remind the team of existing policies) will reduce repeated debates.

By addressing these pain points with the steps above, we aim to **stop the cycle of redoing work**. The development process will become more efficient and predictable – something both the engineering team and I, as PM, desperately need. The end goal is that each piece of work is done once, done well, and integrated smoothly into the product, so we can deliver Gemantria with confidence and not get stuck in endless do-overs.

## Sources:

- Branch health report showing duplicate and stale branches [1](#) [2](#)
  - Integration branch usage and deduplication evidence [5](#)
  - Repeated Ketiv/Qere policy decision in Phase 2 vs Phase 13 [12](#)
  - CI bypass and subsequent fix branches [10](#) [9](#)
- 

[1](#) [2](#) [3](#) [4](#) [5](#) [6](#) [7](#) [8](#) [9](#) [10](#) [13](#) **branch\_health.csv**  
[https://github.com/iog-creator/Gemantria/blob/d01041033f6cf822ccafbd62ad24258e4578a111/archive/\\_reports/branch\\_health\\_20251028\\_175901/branch\\_health.csv](https://github.com/iog-creator/Gemantria/blob/d01041033f6cf822ccafbd62ad24258e4578a111/archive/_reports/branch_health_20251028_175901/branch_health.csv)

[11](#) [12](#) **KETIV\_QERE\_POLICY\_CONFLICT.md**  
[https://github.com/iog-creator/Gemantria/blob/d01041033f6cf822ccafbd62ad24258e4578a111/docs/analysis/KETIV\\_QERE\\_POLICY\\_CONFLICT.md](https://github.com/iog-creator/Gemantria/blob/d01041033f6cf822ccafbd62ad24258e4578a111/docs/analysis/KETIV_QERE_POLICY_CONFLICT.md)

[14](#) [15](#) **branch\_health\_report.py**  
[https://github.com/iog-creator/Gemantria/blob/d01041033f6cf822ccafbd62ad24258e4578a111/scripts/branch\\_health\\_report.py](https://github.com/iog-creator/Gemantria/blob/d01041033f6cf822ccafbd62ad24258e4578a111/scripts/branch_health_report.py)

[16](#) **INDEX.md**  
<https://github.com/iog-creator/Gemantria/blob/d01041033f6cf822ccafbd62ad24258e4578a111/docs/INDEX.md>